

Basic Information

Student name: Yaotian Liu

Project title: RTLRepair-Agent: Verification-Guided Agentic RTL Generation

Repository / notebook link: <https://github.com/ytliu74/rtlrepair-agent>

Configuration location: README.md, .env.example, and examples/.

Section 1. Problem Definition

For an RTL designer, convert a natural-language specification and fixed module interface into synthesizable SystemVerilog. An independent testbench evaluates the output. Success requires compilation and simulation exit codes of zero, an explicit **TEST_PASS**, and no **TEST_FAIL**. Syntax/interface errors, timeouts, mismatches, and missing pass markers are failures. Research question: can compiler and simulator feedback improve an LLM’s functional pass rate over one-shot generation?

Section 2. Motivation and Project Scope

Generated RTL can contain syntax, timing, and functional errors. Compiler diagnostics and simulation mismatches provide objective feedback for an agentic generate–observe–repair workflow. Open-source Icarus Verilog makes evaluation feasible within a semester. Implemented now: one-shot generation and automated verification on a counter and FIFO. Planned: bounded repair and broader evaluation. Synthesis/PPA feedback is a stretch goal; physical design is outside the initial scope.

Section 3. Runnable Baseline

Python 3.10+ calls OpenAI’s pinned `gpt-4.1-2025-04-14` snapshot once per design, cleans/saves the response, compiles with `iverilog -g2012`, and simulates with `vvp`. The model sees only the specification/interface; testbenches are withheld. `src/llm_client.py`, `rtl_generator.py`, `simulator.py`, and `baseline.py` implement the pipeline. `./scripts/run_examples.sh` runs both tasks with two calls total. This control baseline performs no feedback-driven repair or automatic retry.


Section 4. Test Case and Baseline Output

The counter checks synchronous reset, enable/hold, and 8-bit wrap. The more complex FIFO stores sixteen 8-bit words, with registered reads, occupancy/full/empty, ordering, and pointer wrap. At full, simultaneous read/write accepts both; at empty it accepts only the write, with no bypass. An independent shift-queue scoreboard checks directed boundaries and 512 deterministic stimulus cycles.

Actual first GPT-4.1 results on September 6, 2026: both designs compiled, simulated, and emitted **TEST_PASS**. Counter: 853 checks, 1.985 s generation, 329 tokens. FIFO: 744 cycles, 2,233 state checks, 3.478 s generation, 1,099 tokens. Icarus warned that FIFO `unique case` qualities are ignored; compilation exited zero and evaluation passed. Generated RTL was not hand-edited. `generated/{counter,fifo}.sv`, `results/{baseline,fifo}_results.json`, and `artifacts/{baseline,fifo}_output.txt` preserve the actual outputs. Earlier GPT-5.6 counter evidence is archived separately.

Section 5. Reproducibility and Run Instructions

Clone the repository above and enter `rtlrepair-agent`. Install Python 3.10+ and Icarus Verilog (`brew install icarus-verilog` on macOS). Run `python3 -m venv .venv`, `source .venv/bin/activate`, and `python -m pip install -r requirements.txt` (standard library only). Copy `.env.example` to `.env`; fill `OPENAI_API_KEY` and retain `OPENAI_MODEL=gpt-4.1-2025-04-14`. Optional `OPENAI_BASE_URL` defaults to `https://api.openai.com/v1`. Run `./scripts/run_examples.sh`,



```
Compilation: PASS
Simulation: PASS
--- simulation_stdout ---
Checks: 853
TEST_PASS
/Users/henry/repos/CSE598-proposal-RTL/examples/counter/tb.sv:60: $finish called at 2846000 (1ps)
Functional verification: TEST_PASS
Results: results/baseline_results.json
$ ./scripts/run_baseline.sh fifo
[PASS] Python 3.13.1 (requires 3.10+)
[PASS] iverilog found
[PASS] vvp found
[PASS] Python dependencies: standard library only
[PASS] OPENAI_API_KEY configured
[PASS] OPENAI_MODEL configured
[PASS] API URL configuration
Model: gpt-4.1-2025-04-14
Generating RTL (one LLM call; testbench withheld)...
Generation: PASS -> generated/fifo.sv
Compilation: PASS
Simulation: PASS
--- compiler_stderr ---
/Users/henry/repos/CSE598-proposal-RTL/generated/fifo.sv:51: vvp.tgt sorry: Case unique/unique0 qualities are ignored.
--- simulation_stdout ---
Cycles: 744; checks: 2233
Coverage: full_both=4 empty_both=5 both=129 overflow=1 underflow=25 resets=8
TEST_PASS
/Users/henry/repos/CSE598-proposal-RTL/examples/fifo/tb.sv:116: $finish called at 7446000 (1ps)
Functional verification: TEST_PASS
Results: results/fifo_results.json
=== One-shot suite (one API call per design) ===
Counter: PASS
FIFO: PASS

CSE598-proposal-RTL on main [!?] via v3.11.4 via base took 6s
```

Figure 1: Genuine Terminal capture: GPT-4.1 FIFO command, verification output, and both task results.

or `./scripts/run_baseline.sh fifo` for FIFO alone. Each `examples/<task>/` contains `spec.txt` and `tb.sv`; outputs are listed above. API model access is required. README supplies Linux setup and offline verification; 20 tests pass.

Section 6. Initial Evaluation Plan

Compare one-shot generation with bounded verification-guided repair using the same model, initial prompt, tasks, and independent testbenches over repeated trials. Primary metric: functional pass rate. Also measure compile rate, iterations, calls, tokens, latency, and measurable cost. Include budget-matched independent resampling. Expand curated examples before considering VerilogEval or RTLLM. Retain all failed runs.

Section 7. Limitations and Next Steps

Two tasks cannot establish generality or a model ranking; state comparisons are not separate design tasks. Simulation covers only tested behavior and does not prove synthesizability. Model nondeterminism, API availability, and testbench coverage limit reliability. Next: implement diagnosis and bounded repair, expand evaluation, and optionally incorporate synthesis/PPA feedback.